


Learning representations by back-propagating errors

Rumelhart, David E., Geoffrey E. Hinton,
and Ronald J. Williams.

수학과 오서영

Neural Network



Back-propagation

If input units are directly connected to the output units

- > easy to find learning rules
(iteratively adjust the relative strengths)
- > progressively reduce the difference
between the actual and desired output

Hidden units

- > Learning procedure decides under what
circumstances the hidden units should be active in order to achieve the
desired input-output
- > Hidden units are going to learn to represent some features of input
domain
- > **Back propagation**

Feed Forward

Input x_j

- > input of higher layer will be a weighted sum of the outputs of all of the lower layer units that feed into it
- > i feeds into j

Output y_i

Adjusted by a weight on the link between i and j w_{ij}

$$x_j = \sum y_i w_{ij}$$

Feed Forward

Output y_j

-> **Activation function** in this paper is the **sigmoid function**

$$y_j = \frac{1}{1 + e^{-x_j}}$$

-> By repeating this procedure, starting with the **input layer**, we can feed-forward through the network layers and arrive at a set of **outputs for the output layer**

Error Function

The aim is to find a set of weights that ensure that the output vector is the same as the desired output vector

- > computed output value of j in output layer y_j
- > its desired state d_j

$$E = \frac{1}{2} \sum_j (y_j - d_j)^2$$

- > minimize
- > need to adjust the weights

Error Function

To minimize E by gradient descent it is necessary to compute the **partial derivative** of E with respect to each weight in the network

$$\rightarrow \frac{\partial E}{\partial w} \quad \rightarrow \frac{\partial E}{\partial y_j} \text{ (E and } y \text{ are directly related)} \quad \rightarrow \frac{\partial y_j}{\partial x_j}$$

$$\frac{\partial}{\partial y_j} \frac{1}{2} \sum_j (y_j - d_j)^2 = y_j - d_j \quad y_j = \text{sigmoid}(x_j)$$

$$\frac{\partial y_j}{\partial x_j} = \text{sigmoid}(x_j) * (1 - \text{sigmoid}(x_j))$$

Error Function

$$x_j = \sum y_i w_{ij}$$

$$\rightarrow \frac{\partial x_j}{\partial w_{ij}} = y_i$$

We can chain them together to figure out how the error changes with the weights

$$\frac{\partial E}{\partial w} = \frac{\partial E}{\partial y} * \frac{\partial y}{\partial x} * \frac{\partial x}{\partial w}$$

Gradient Descent

Simplest form of **gradient descent** would be to change each weight by an amount proportional to the accumulated gradient

$$\Delta w = -a * \frac{\partial E}{\partial w}$$

Code implementation

```
display(Math(r'x_j = \sum_{i=0} y_i w_{ij}'))
display(Math(r'y_j = \frac{1}{1+e^{-x_j}}'))
display(Math(r'E = \frac{1}{2} \sum_j (y_j - d_j)^2'))
display(Math(r'\frac{\partial E}{\partial y_j} = y_j - d_j'))
display(Math(r'\frac{\partial y_j}{\partial x_j} = y_j * (1 - y_j)'))
display(Math(r'\frac{\partial x_j}{\partial w_{ij}} = y_i'))
```

$$x_j = \sum_{i=0} y_i w_{ij}$$

$$y_j = \frac{1}{1 + e^{-x_j}}$$

$$E = \frac{1}{2} \sum_j (y_j - d_j)^2$$

$$\frac{\partial E}{\partial y_j} = y_j - d_j$$

$$\frac{\partial y_j}{\partial x_j} = y_j * (1 - y_j)$$

$$\frac{\partial x_j}{\partial w_{ij}} = y_i$$

```
def sigmoid(xj):  
  
    s = 1/(1+np.exp(-xj))  
  
    return s
```

Code implementation

```
: def propagate(w, yi, yj, dj):  
  
    m = yi.shape[1]  
  
    yj = sigmoid(np.dot(w.T, yi))           # compute activation  
    cost = 1/m*(np.sum(yj-dj))             # compute cost  
  
    dw = 1/m*np.dot((np.dot(yi, ((yj-dj).T))), yj*(1-yj))  
  
    cost = np.squeeze(cost)  
  
    return dw, cost
```

Code implementation

```
def optimize(w, yi, dj, num_iterations, learning_rate):  
    costs = []  
  
    for i in range(num_iterations):  
  
        dw, cost = propagate(w, yi, yj, dj)  
  
        w = w - learning_rate*dw  
  
        if i % 1 == 0:  
            costs.append(cost)  
  
        if i % 1 == 0:  
            print ("Cost after iteration %i: %f" %(i, cost))  
  
    return w, dw, yj, costs
```

Code implementation

```
w, yi, dj = np.array([[1.],[2.]]), np.array([[1,2,-1],[3,4,-3.2]]), np.array([[1,0,1]])
```

```
w, dw, yj, costs = optimize(w, yi, dj, num_iterations= 10, learning_rate = 0.1)
```

```
print ("w = " + str(w) + '\n')
print ("dw = " + str(dw) + '\n')
print("prediction = " + str(yj))
```

```
Cost after iteration 0: -0.000115
```

```
Cost after iteration 1: -0.000346
```

```
Cost after iteration 2: -0.000346
```

```
Cost after iteration 3: -0.000347
```

```
Cost after iteration 4: -0.000347
```

```
Cost after iteration 5: -0.000347
```

```
Cost after iteration 6: -0.000348
```

```
Cost after iteration 7: -0.000348
```

```
Cost after iteration 8: -0.000349
```

```
Cost after iteration 9: -0.000349
```

```
w = [[0.997439  0.99987199 0.99828173]
      [1.99385446 1.99969281 1.99587673]]
```

```
dw = [[0.00275869 0.00013821 0.0018515 ]
      [0.00661992 0.00033165 0.00444297]]
```

```
prediction = [[1 0 1]]
```