

Reducing the Dimensionality of Data with Neural Networks

Hinton, Geoffrey E.,
and Ruslan R. Salakhutdinov.

Overview

Dimensionality reduction

- classification, visualization, communication, storage of high-dimensional data

Goal : use a deep neural to reduce the dimensionality of an input.

-> the reduction techniques described are compared to principal component analysis (PCA) and logistic PCA

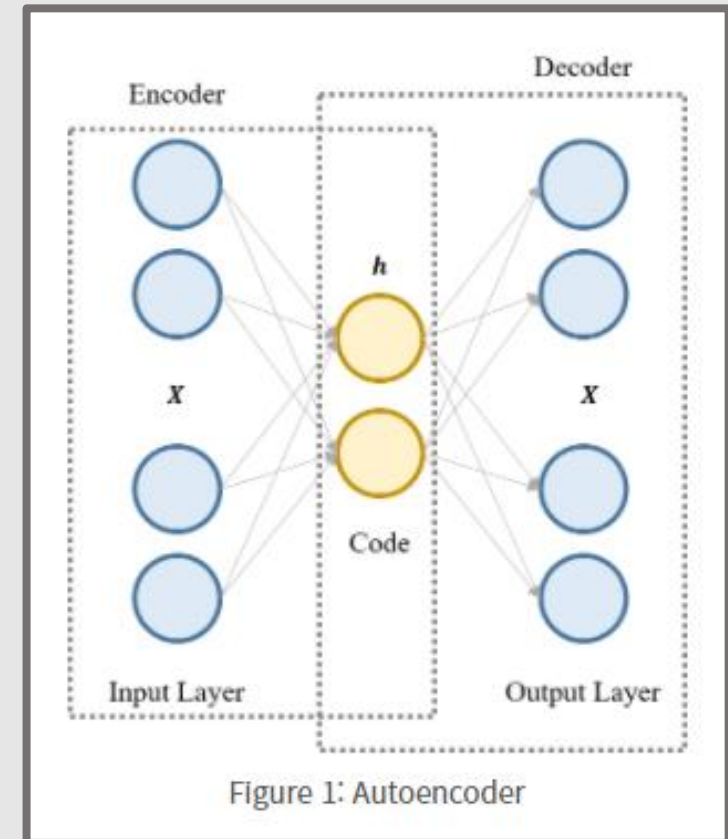
PCA

: Find the directions of greatest variance in the data set and
Represent each data point by its coordinates along each of these directions.

Overview

Autoencoder

- Nonlinear generalization of PCA
- Encoder transform the high-dimensional data into a low-dimensional code
- Decoder recover the data from the code
- Starts with random weights in the two networks
- Trained by minimizing the discrepancy between the original data and its reconstruction.
- Gradients are obtained by the chain rule to back-propagate error from the decoder network to encoder network.



Overview

It is difficult to optimize multilayer autoencoder

1. With large initial weights
: autoencoders typically find poor local minima
2. With small initial weights
: the gradients in the early layers are tiny, making it infeasible to train autoencoders with many hidden layers
3. If initial weights are close to a good solution
: gradient descent works well
-> **But finding such initial weights is very difficult**

Pretraining procedure for binary data, generalize it to real-valued data, and show that it works well for a variety of data sets.

-> based on **Restricted Boltzmann Machine (RBM)**

Methods



Restricted Boltzmann Machine (RBM)

The input data correspond to **visible** units of the RBM and the feature detectors correspond to hidden units.

A joint configuration (w, h) of the visible and hidden units has an **energy**

The network assigns a probability to every possible data via this energy function

$$E\{h_0x_0 - h_{\infty}x_{\infty}\}$$

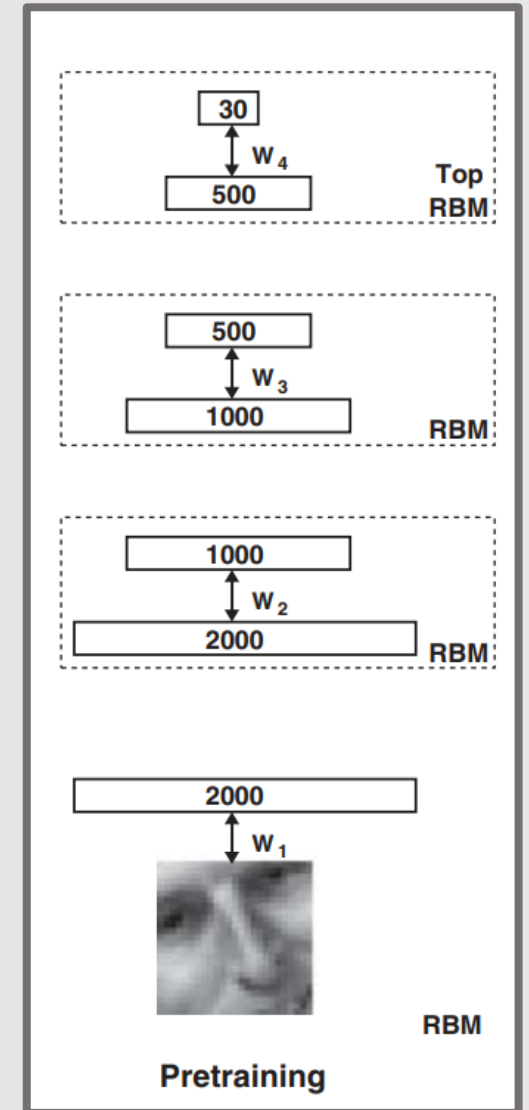


Methods

Training a large neural network through greedy layer-wise pretraining

He took an input (50*50 images), and trained a RBM to reduce the feature space of the image.

Each set of weights are learned individually, so in the first step, W_1 is learned and once W_1 is learned, the data is then all restored as the result of the first transfer and W_2 is learned, and so on.



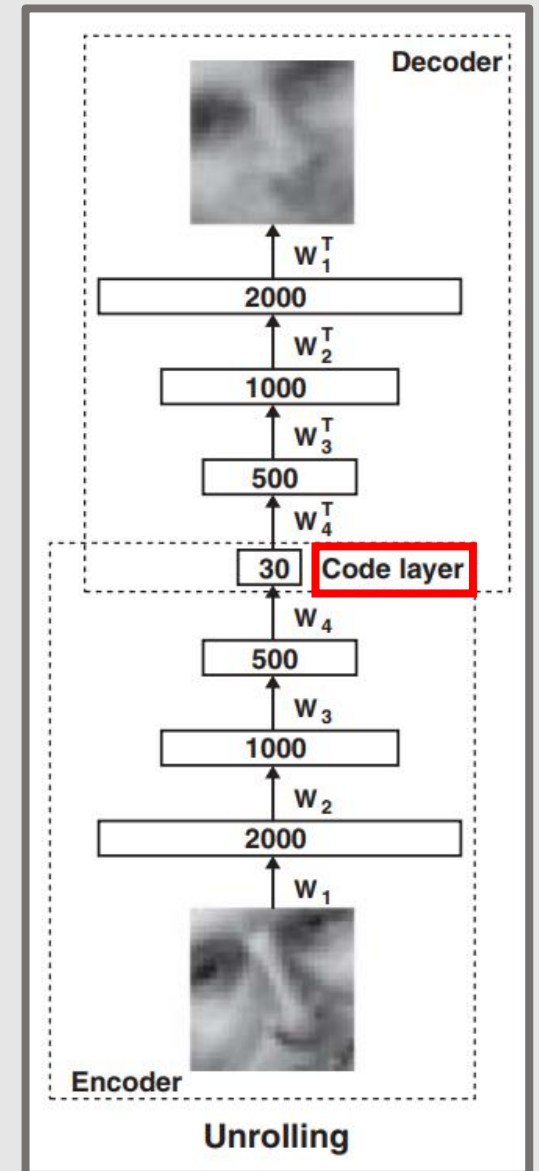
Methods

RBM is generative models, so the weights are bidirectional.
-> activation in layer two can be used to generate a corresponding set of activations in layer one.

The weights between the layers are effectively re-used to map the data back to its original dimensionality.

-> the input weights for a NN that can be trained with backpropagation to fine-tune the weights using the input weights as the outputs to the network

Code layer : used for classification, regression, clustering, etc.



Results

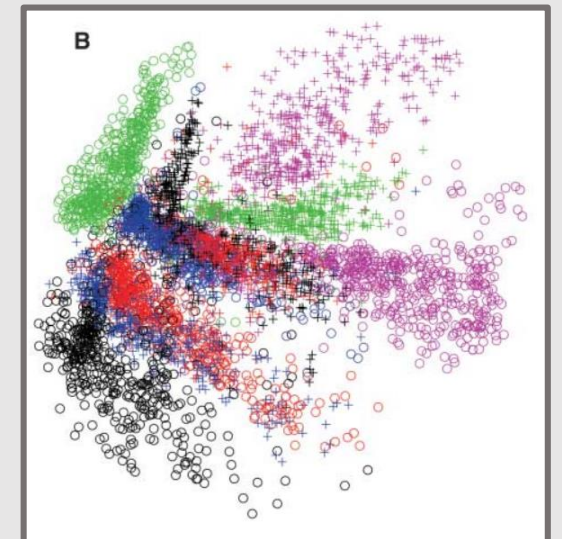

MNIST

A : 2 dimensional results of standard PCA

B : results of a deep neural network

-> autoencoder produces a much more distinct separation of digits when projected into 2 dimensional space than PCA.

-> because the autoencoder has the ability to learn non-linear features since each layers has non-linear activation units and thus can learn a more complex representation of the input.



Analysis

Initialize weights using RBM & gradient descent

- > learning ↑
- > deeplearning ↑

Deep Belief Network (DBN)

: Restricted Boltzmann machine with connection only between layer and layer

- > learning ↑

General NN : Learn weights from upper to lower layers

DBN : Learn weights from lower to upper layers